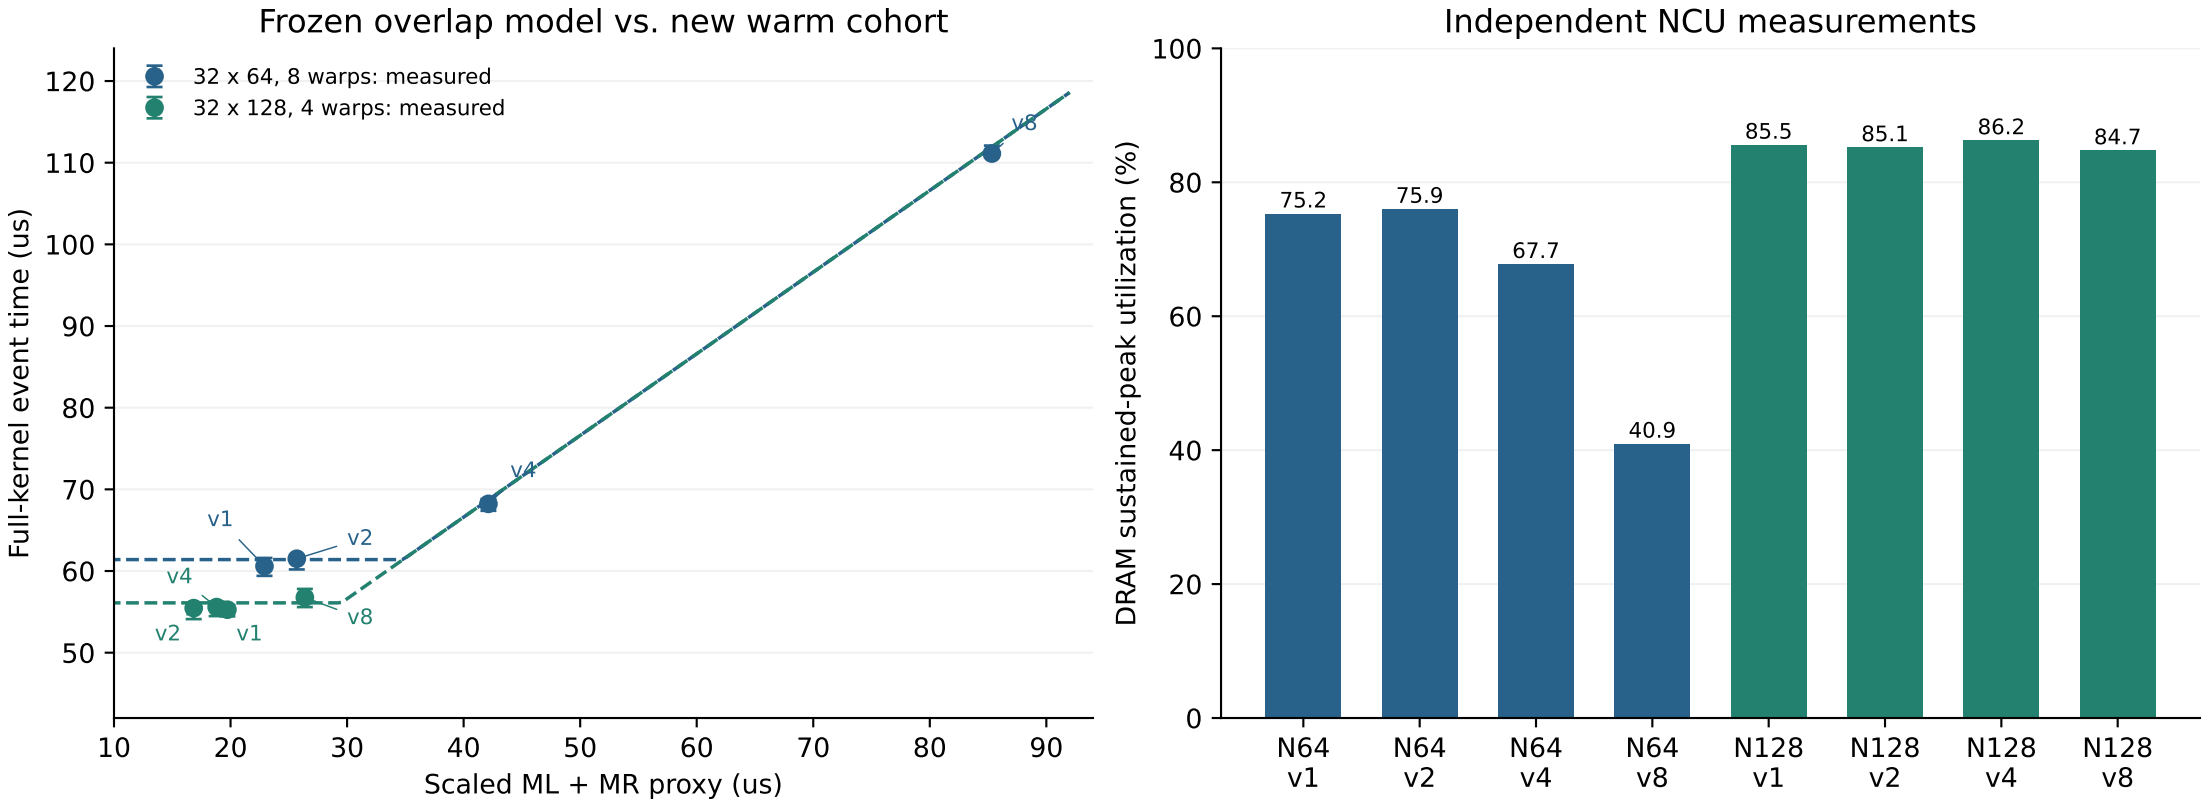
Frozen overlap model vs. new warm cohort
32 x 64, 8 warps: measured
32 x 128, 4 warps: measured
Full-kernel event time (us)
Scaled ML + MR proxy (us)
v1
v2
v4
v8
Independent NCU measurements
DRAM sustained-peak utilization (%)
75.2
75.9
67.7
40.9
85.5
85.1
86.2
84.7
N64 v1
N64 v2
N64 v4
N64 v8
N128 v1
N128 v2
N128 v4
N128 v8